

Lecture 05: The Process Abstraction

Vivek Kumar

Computer Science and Engineering

IIIT Delhi

vivekk@iiitd.ac.in

Last Lecture: Program Execution (1/2)

```
L1: int g=0;
```

```
L2: void main() {
```

```
L3:   int *a = (int*) malloc(4);
```

```
L4:   char *b = "Hello World";
```

```
L5:   foo(a);
```


```
L6:   g=*a;
```

```
L7: }
```

```
L8: void foo(int* b) {
```

```
L9:   *b = 20;
```

```
L10: }
```




Last Lecture: Program Execution (2/2)

L1: `int g=0;`

L2: `void main() {`
 L3: `int *a = (int*) malloc(4);`
 L4: `char *b = "Hello World";`
 L5: `foo(a);`
 L6: `g=*a;`
 L7: `}`

L8: `void foo(int* b) {`
 L9: `*b = 20;`
 L10: `}`



The Story So Far...

- A **cook** (OS) manages a **kitchen** (*computer*)
 - Lecture 01
- The job of a cook is to do **cooking** (*execute binaries*) by reading **recipes** (*executable file*) in a **standard format** (*ELF*)
 - Lecture 02
- The cook has received recipe of a dish in **different files** (*relocatable files*). It will have to use a **stapler** (*linker*) to stitch these files to make a final recipe
 - Lecture 03
- The cook has to **store** (*load*) the **steps** (*instructions*) from the recipe in his memory to start cooking. The cook then execute instructions **one by one** (*call stack*) for cooking
 - Lecture 04
- The kitchen has gained a lot of popularity and have started receiving lot of recipes for cooking. The cook has hired several **helpers** (*processes*), where one helper cooks one recipe
 - **Today's point of discussion!**

Today's Class

- The process abstraction
- Quiz-1 (Lectures 02-04)

All Executables are ELF

```
iiitd@possum:~$ file /bin/cat
/bin/cat: ELF 64-bit LSB shared object, x86-64, version 1 (SYSV), dynamically
linked, interpreter /lib64/ld-linux-x86-64.so.2, for GNU/Linux 3.2.0, Build
ID[sha1]=747e524bc20d33ce25ed4aea108e3025e5c3b78f, stripped
iiitd@possum:~$ file /bin/ls
/bin/ls: ELF 64-bit LSB shared object, x86-64, version 1 (SYSV), dynamically
linked, interpreter /lib64/ld-linux-x86-64.so.2, for GNU/Linux 3.2.0, Build
ID[sha1]=9567f9a28e66f4d7ec4baf31cfbf68d0410f0ae6, stripped
iiitd@possum:~$ file /bin/mkdir
/bin/mkdir: ELF 64-bit LSB shared object, x86-64, version 1 (SYSV), dynamica
lly linked, interpreter /lib64/ld-linux-x86-64.so.2, for GNU/Linux 3.2.0, Bu
ildID[sha1]=6c825e92c5eae304845d070ed34749495d67c566, stripped
```



The Process

```
iiitd@possum:~$ vi fib.c
iiitd@possum:~$ gcc fib.c
iiitd@possum:~$ ./a.out
Fib(40) = 102334155
```

PID	USER	PRI	NI	VIRT	RES	SHR	S	CPU%	MEM%	TIME+	Command
2034	iiitd	20	0	4384	820	756	R	100.	0.0	0:17.48	./a.out

- Whenever we type and enter a command on shell, the command prompt is returned back only after that command has completed its execution
 - How this execution happens?
- A program under execution is called a process
 - It is again all about abstraction. First we discussed about ELF to have an abstraction and presenting a unified view of all object files for the loader. Now, we are going one level down where the OS needs to have an abstract representation of any kind of program under execution
 - Process is a collection of resources

There are Several Processes

```

iiitd@possum:~$ for i in {1..20}
> do
> ./a.out &
> done
[1] 2062
[2] 2063
[3] 2064
[4] 2065
[5] 2066
[6] 2067
[7] 2068
[8] 2069
[9] 2070
[10] 2071
[11] 2072
[12] 2073
[13] 2074
[14] 2075
[15] 2076
[16] 2077
[17] 2078
[18] 2079
[19] 2080
[20] 2081

```

```

1  |||||100.0% Tasks: 104, 202 thr; 4
2  |||||100.0% Load average: 2.36
3  |||||100.0% Uptime: 00:09:59
4  |||||100.0%
Mem ||||| 470M/3.75G
Swp ||||| 0K/2.00G

```

PID	USER	PRI	NI	VIRT	RES	SHR	S	CPU%	MEM%	TIME+	Command
2074	iiitd	20	0	4384	804	740	R	20.5	0.0	0:06.97	./a.out
2076	iiitd	20	0	4384	768	704	R	20.5	0.0	0:06.96	./a.out
2080	iiitd	20	0	4384	792	728	R	20.5	0.0	0:06.96	./a.out
2068	iiitd	20	0	4384	712	648	R	20.5	0.0	0:06.96	./a.out
2066	iiitd	20	0	4384	820	756	R	19.9	0.0	0:06.97	./a.out
2064	iiitd	20	0	4384	708	648	R	19.9	0.0	0:06.97	./a.out
2072	iiitd	20	0	4384	756	692	R	19.9	0.0	0:06.96	./a.out
2063	iiitd	20	0	4384	800	740	R	19.9	0.0	0:06.96	./a.out
2069	iiitd	20	0	4384	720	656	R	19.9	0.0	0:06.96	./a.out
2078	iiitd	20	0	4384	864	800	R	19.9	0.0	0:06.98	./a.out
2070	iiitd	20	0	4384	804	740	R	19.9	0.0	0:06.98	./a.out
2075	iiitd	20	0	4384	752	692	R	19.9	0.0	0:06.96	./a.out
2073	iiitd	20	0	4384	764	700	R	19.9	0.0	0:06.98	./a.out
2065	iiitd	20	0	4384	804	740	R	19.9	0.0	0:06.98	./a.out
2077	iiitd	20	0	4384	712	648	R	19.9	0.0	0:06.93	./a.out
2071	iiitd	20	0	4384	752	692	R	19.9	0.0	0:06.94	./a.out
2081	iiitd	20	0	4384	756	692	R	19.9	0.0	0:06.93	./a.out
2079	iiitd	20	0	4384	720	656	R	19.9	0.0	0:06.94	./a.out
2067	iiitd	20	0	4384	716	656	R	19.9	0.0	0:06.93	./a.out
2062	iiitd	20	0	4384	800	740	R	19.9	0.0	0:06.94	./a.out
2021	iiitd	20	0	33784	4536	3612	R	0.0	0.1	0:00.68	htop

```

iiitd@possum:~$ cat /proc/sys/kernel/pid_max
32768

```

```

iiitd@possum:~$ lscpu
Architecture:          x86_64
CPU op-mode(s):        32-bit, 64-bit
Byte Order:            Little Endian
CPU(s):                 4
On-line CPU(s) list:   0-3
Thread(s) per core:    1
Core(s) per socket:    4

```

- Note the “R”(running) state of the program in “S”(state) column in htop output
- There are only four CPUs in this system, but how is the OS able to run 20 a.out simultaneously?
- Virtualization of the CPU!
 - Any limit?

Process Keep Changing its State

```

iiitd@possum:~$ ./a.out
Enter input..
40
Fib(40) = 102334155
Enter input..
50

```


PID	USER	PRI	NI	VIRT	RES	SHR	S	CPU%	MEM%	TIME+	Command
2327	iiitd	20	0	4516	760	700	S	0.0	0.0	1:15.50	./a.out

PID	USER	PRI	NI	VIRT	RES	SHR	S	CPU%	MEM%	TIME+	Command
2327	iiitd	20	0	4516	760	700	R	99.5	0.0	1:21.13	./a.out

PID	USER	PRI	NI	VIRT	RES	SHR	S	CPU%	MEM%	TIME+	Command
2327	iiitd	20	0	4516	760	700	S	0.0	0.0	1:15.50	./a.out

- There are two distinct phases in the above program
 - Asking the user for the input
 - The process representing the a.out is moved into “S” state
 - It got kicked out of the CPU
 - Calculating the Fibonacci number of that input
 - The process representing the a.out is moved into “R” state
 - It got its CPU back!
 - **Why not simply leave the process occupying the CPU?**

Process States



- As a process executes, it changes **state**
 - **New**: The process is being created
 - Loader loading the executable
 - **Running**: Instructions are being executed
 - **Waiting**: The process is waiting for some event to occur
 - **Ready**: The process is waiting to be assigned to a processor
 - **Terminated**: The process has finished execution

Process State: Example

Time	Process ₀	Process ₁	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	Process ₀ initiates I/O
4	Blocked	Running	Process ₀ is blocked,
5	Blocked	Running	so Process ₁ runs
6	Blocked	Running	
7	Ready	Running	I/O done
8	Ready	Running	Process ₁ now done
9	Running	–	
10	Running	–	Process ₀ now done

Figure 4.4: Tracing Process State: CPU and I/O

What Constitutes a Process

- There are so many processes at any given time
 - Each process has its state
 - A running process will be using some CPU registers
 - We push and pop the registers on call stack during the course of execution
 - Process requires memory for execution (stack and heap)
 - Process have several open file handles
- For process management, the OS uses a data-structure called as “process descriptor” or “process control block” to keep track of every process’s progress and usage of the computer’s available resources

Process Descriptor (a.k.a PCB)

- Kernel maintains detail on each process inside a structure of type **task_struct**
- You can see it yourself
 - <https://raw.githubusercontent.com/torvalds/linux/master/include/linux/sched.h>

```

struct task_struct {
    unsigned int          __state;
    void                 *stack;
    int                   prio;
    int                   nr_cpus_allowed;
    struct mm_struct      *mm;
    struct mm_struct      *active_mm;
    pid_t                 pid;
    struct list_head      children;
    struct list_head      sibling;
    struct task_struct    *group_leader;

    /*
     * executable name, excluding path.
     *
     * - normally initialized setup_new_exec()
     * - access it with [gs]et_task_comm()
     * - lock it with task_lock()
     */
    char                   comm[TASK_COMM_LEN];

    struct nameidata      *nameidata;
    /* Filesystem information: */
    struct fs_struct      *fs;

    /* Open file information: */
    struct files_struct    *files;

```

- Extracts from the struct `task_struct`
 - There are many more members!
- Registers are saved on stack
- Linked list of PCB to dynamically add a new process information

```

iiitd@possum:~$ ps -u iiitd --forest
  PID TTY          TIME CMD
 23005 ?            00:00:00 sshd
 23006 pts/1        00:00:00 \_  bash
 23062 pts/1        00:00:01 \_  a.out
 22906 ?            00:00:00 sshd
 22908 pts/0        00:00:00 \_  bash
 23063 pts/0        00:00:00 \_  ps
 22807 ?            00:00:00 systemd
 22808 ?            00:00:00 \_  (sd-pam)

```

Next Lecture

- System calls
- Assignment-1 will be released today with a deadline of 2nd September